

# cereal

February 21, 2020

## 1 HW 1 - Visualizing Cereals Dataset

Elmer Camargo

```
[1]: import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
from plotnine import *
from plotnine.data import mtcars
import seaborn as sns
from scipy import stats

get_ipython().run_line_magic('matplotlib', 'inline')

def r2(x, y):
    return stats.pearsonr(x, y)[0] ** 2
```

```
[2]: cereal = pd.read_csv("data\cereal.csv")
cereal.head()
```

```
[2]:
```

	name	mfr	type	calories	protein	fat	sodium	fiber	\
0	100%_Bran	N	C	70	4	1	130	10.0	
1	100%_Natural_Bran	Q	C	120	3	5	15	2.0	
2	All-Bran	K	C	70	4	1	260	9.0	
3	All-Bran_with_Extra_Fiber	K	C	50	4	0	140	14.0	
4	Almond_Delight	R	C	110	2	2	200	1.0	

	carbo	sugars	potass	vitamins	shelf	weight	cups	rating
0	5.0	6.0	280.0	25	3	1.0	0.33	68.402973
1	8.0	8.0	135.0	0	3	1.0	1.00	33.983679
2	7.0	5.0	320.0	25	3	1.0	0.33	59.425505
3	8.0	0.0	330.0	25	3	1.0	0.50	93.704912
4	14.0	8.0	NaN	25	3	1.0	0.75	34.384843

```
[3]: cereal.info()
cereal.isnull().sum()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 77 entries, 0 to 76
Data columns (total 16 columns):
name          77 non-null object
mfr           77 non-null object
type          77 non-null object
calories      77 non-null int64
protein       77 non-null int64
fat           77 non-null int64
sodium        77 non-null int64
fiber         77 non-null float64
carbo         76 non-null float64
sugars        76 non-null float64
potass        75 non-null float64
vitamins      77 non-null int64
shelf         77 non-null int64
weight        77 non-null float64
cups          77 non-null float64
rating        77 non-null float64
dtypes: float64(7), int64(6), object(3)
memory usage: 9.8+ KB

```

```

[3]: name          0
     mfr           0
     type          0
     calories      0
     protein       0
     fat           0
     sodium        0
     fiber         0
     carbo         1
     sugars        1
     potass        2
     vitamins      0
     shelf         0
     weight        0
     cups          0
     rating        0
     dtype: int64

```

```
[4]: cereal[cereal.isna().any(axis=1)]
```

```

[4]:
      name mfr type  calories  protein  fat  sodium  fiber  \
4   Almond_Delight  R   C      110         2    2     200    1.0
20 Cream_of_Wheat_(Quick)  N   H      100         3    0      80    1.0
57   Quaker_Oatmeal  Q   H      100         5    2       0    2.7

```

	carbo	sugars	potass	vitamins	shelf	weight	cups	rating
4	14.0	8.0	NaN	25	3	1.0	0.75	34.384843
20	21.0	0.0	NaN	0	2	1.0	1.00	64.533816
57	NaN	NaN	110.0	0	1	1.0	0.67	50.828392

```
[5]: cereal = cereal.fillna(0)
cereal.describe()
```

```
[5]:
```

	calories	protein	fat	sodium	fiber	carbo \
count	77.000000	77.000000	77.000000	77.000000	77.000000	77.000000
mean	106.883117	2.545455	1.012987	159.675325	2.151948	14.610390
std	19.484119	1.094790	1.006473	83.832295	2.383364	4.232257
min	50.000000	1.000000	0.000000	0.000000	0.000000	0.000000
25%	100.000000	2.000000	0.000000	130.000000	1.000000	12.000000
50%	110.000000	3.000000	1.000000	180.000000	2.000000	14.000000
75%	110.000000	3.000000	2.000000	210.000000	3.000000	17.000000
max	160.000000	6.000000	5.000000	320.000000	14.000000	23.000000

	sugars	potass	vitamins	shelf	weight	cups \
count	77.000000	77.000000	77.000000	77.000000	77.000000	77.000000
mean	6.935065	96.103896	28.246753	2.207792	1.029610	0.821039
std	4.422840	71.251147	22.342523	0.832524	0.150477	0.232716
min	0.000000	0.000000	0.000000	1.000000	0.500000	0.250000
25%	3.000000	40.000000	25.000000	1.000000	1.000000	0.670000
50%	7.000000	90.000000	25.000000	2.000000	1.000000	0.750000
75%	11.000000	120.000000	25.000000	3.000000	1.000000	1.000000
max	15.000000	330.000000	100.000000	3.000000	1.500000	1.500000


	rating
count	77.000000
mean	42.665705
std	14.047289
min	18.042851
25%	33.174094
50%	40.400208
75%	50.828392
max	93.704912

```
[6]: #g = sns.pairplot(cereal)
```


### 1.0.1 Relationship Between Fiber and Sodium

It appears that most cereals with low fiber have mid to high sodium values (150 - 300mg) as well as typically more calories, however there are observations of cereals with both low fiber and low sodium as seen on the left half of the scatterplot. Additionally, it does not seem that sodium and fiber have an immediate relationship based off of the regression plot and r2.

```
[7]: g = sns.scatterplot(data = cereal, x = 'sodium', y = 'fiber', hue = "mfr", size_
    ↳ "calories",
    sizes=(20, 200), hue_norm=(0, 7), legend = False)
```



```
[8]: #g = sns.regplot(data = cereal, x = 'sodium', y = 'fiber')
g = sns.jointplot(data = cereal, x = 'sodium', y = 'fiber', kind="reg",
    ↳ stat_func=r2)
```




```
[9]: cereal_bar = cereal.groupby("shelf").agg({"sugars": ["mean", "std", "min"]})
cereal_bar.columns = ["avg_sugars_g", "std", "min"]
cereal_bar["shelf"] = cereal_bar.index #how you grab row names
cereal_bar
```

```
[9]:      avg_sugars_g      std  min  shelf
shelf
1          4.850000  4.510514  0.0      1
2          9.619048  4.128876  0.0      2
3          6.527778  3.835817  0.0      3
```

### 1.0.2 Good cereals are usually on the 2nd or 3rd shelf

couldnt figure out how to fix legend

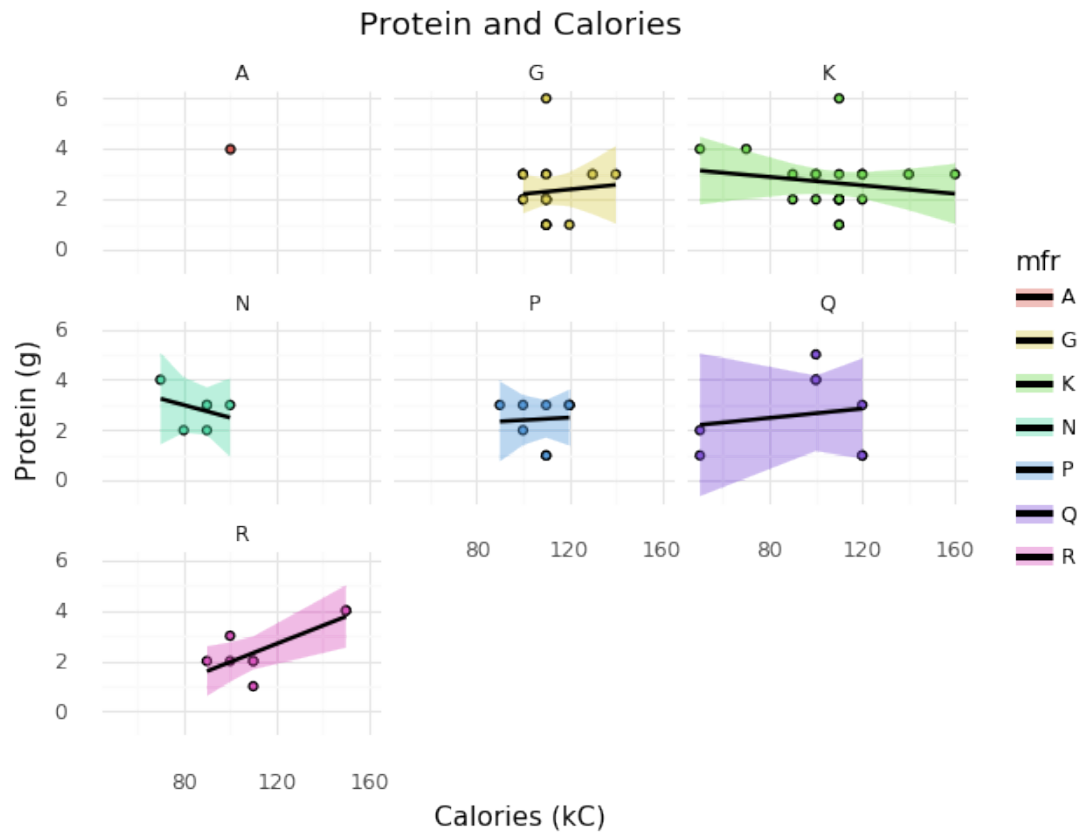
```
[10]: (ggplot(cereal_bar, aes( x = "shelf", y = "avg_sugars_g", fill = "shelf"))
+ geom_bar(stat = "identity"))
```



```
[10]: <ggplot: (-9223371927783853816)>
```

## 1.1 Protein and Calories by Manufacturer

```
[11]: (ggplot(cereal, aes( x = 'calories', y = 'protein'))
+ geom_point(aes(fill = 'mfr'), show_legend = False)
+ stat_smooth(aes(fill = 'mfr'), method = 'lm')
+ facet_wrap('~mfr')
+ labs(x = "Calories (kC)", y = "Protein (g)", title = "Protein and
↳ Calories")
+ theme_minimal())
```



[11]: <ggplot: (-9223371927783454356)>